

investigate-a-dataset-TMDb

January 17, 2019

1 Project: TMDb Movie Data Analysis

1.1 Table of Contents

Introduction

Data Wrangling

Exploratory Data Analysis

Conclusions

Introduction

In this report we work on a data set from TMDb. The data set contains information about 10,000 movies collected from The Movie Database (TMDb), including user ratings, genres, movie budget and movie revenue. In particular we want to understand: 1. Which genres were most popular between 1996 and 2015? 2. What kind of properties are associated with movies that have high revenues?

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

Data Wrangling

1.1.1 Reading Data

```
In [2]: df = pd.read_csv('tmdb-movies.csv')
df.head()
```

```
Out[2]:
```

	id	imdb_id	popularity	budget	revenue	\
0	135397	tt0369610	32.985763	150000000	1513528810	
1	76341	tt1392190	28.419936	150000000	378436354	
2	262500	tt2908446	13.112507	110000000	295238201	
3	140607	tt2488496	11.173104	200000000	2068178225	
4	168259	tt2820852	9.335014	190000000	1506249360	

	original_title	\
0	Jurassic World	
1	Mad Max: Fury Road	
2	Insurgent	

3 Star Wars: The Force Awakens
 4 Furious 7

cast \
 0 Chris Pratt|Bryce Dallas Howard|Irrfan Khan|Vi...
 1 Tom Hardy|Charlize Theron|Hugh Keays-Byrne|Nic...
 2 Shailene Woodley|Theo James|Kate Winslet|Ansel...
 3 Harrison Ford|Mark Hamill|Carrie Fisher|Adam D...
 4 Vin Diesel|Paul Walker|Jason Statham|Michelle ...

homepage director \
 0 http://www.jurassicworld.com/ Colin Trevorrow
 1 http://www.madmaxmovie.com/ George Miller
 2 http://www.thedivergentseries.movie/#insurgent Robert Schwentke
 3 http://www.starwars.com/films/star-wars-episod... J.J. Abrams
 4 http://www.furious7.com/ James Wan

tagline ... \
 0 The park is open. ...
 1 What a Lovely Day. ...
 2 One Choice Can Destroy You ...
 3 Every generation has a story. ...
 4 Vengeance Hits Home ...

overview runtime \
 0 Twenty-two years after the events of Jurassic ... 124
 1 An apocalyptic story set in the furthest reach... 120
 2 Beatrice Prior must confront her inner demons ... 119
 3 Thirty years after defeating the Galactic Empi... 136
 4 Deckard Shaw seeks revenge against Dominic Tor... 137

genres \
 0 Action|Adventure|Science Fiction|Thriller
 1 Action|Adventure|Science Fiction|Thriller
 2 Adventure|Science Fiction|Thriller
 3 Action|Adventure|Science Fiction|Fantasy
 4 Action|Crime|Thriller

production_companies release_date vote_count \
 0 Universal Studios|Amblin Entertainment|Legenda... 6/9/15 5562
 1 Village Roadshow Pictures|Kennedy Miller Produ... 5/13/15 6185
 2 Summit Entertainment|Mandeville Films|Red Wago... 3/18/15 2480
 3 Lucasfilm|Truenorth Productions|Bad Robot 12/15/15 5292
 4 Universal Pictures|Original Film|Media Rights ... 4/1/15 2947

vote_average release_year budget_adj revenue_adj
 0 6.5 2015 1.379999e+08 1.392446e+09
 1 7.1 2015 1.379999e+08 3.481613e+08

2	6.3	2015	1.012000e+08	2.716190e+08
3	7.5	2015	1.839999e+08	1.902723e+09
4	7.3	2015	1.747999e+08	1.385749e+09

[5 rows x 21 columns]

In [3]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
id                10866 non-null int64
imdb_id           10856 non-null object
popularity        10866 non-null float64
budget            10866 non-null int64
revenue           10866 non-null int64
original_title    10866 non-null object
cast              10790 non-null object
homepage          2936 non-null object
director          10822 non-null object
tagline           8042 non-null object
keywords          9373 non-null object
overview          10862 non-null object
runtime           10866 non-null int64
genres            10843 non-null object
production_companies 9836 non-null object
release_date      10866 non-null object
vote_count        10866 non-null int64
vote_average      10866 non-null float64
release_year      10866 non-null int64
budget_adj        10866 non-null float64
revenue_adj       10866 non-null float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB
```

1.1.2 Unuseful Feilds

1. There are some fields which are not relevnat in our analysis like "imdb_id", "original_title", "homepage", "tagline", "keywords" and "overview". We'll drop them in the data cleaning step.
2. The last two fields of "budget_adj" and "revenue_adj" will be used instead of "budget" and "revenue" because they take into account the inflation over the years.
3. 'cast', 'director' and 'production_companies' are not going to be used in our analysis, because there are so many differnt values in them, making it hard to categorize and associate them with other fields.

- There are two fields for release date: "release_date" and "release_year". "release_date" has more detail information which we don't need for our analysis, so we keep "release_year" and drop "release_date".

1.1.3 Missing Values

There are some missing values in 'imdb_id', 'cast', 'homepage', 'director', 'tagline', 'keywords', 'genres' and 'production_companies'. In the previous section, we decided to drop all these fields except "genres". So we just need to take action on "genres" missing values.

```
In [4]: df.describe()
```

```
Out [4]:
```

	id	popularity	budget	revenue	runtime \
count	10866.000000	10866.000000	1.086600e+04	1.086600e+04	10866.000000
mean	66064.177434	0.646441	1.462570e+07	3.982332e+07	102.070863
std	92130.136561	1.000185	3.091321e+07	1.170035e+08	31.381405
min	5.000000	0.000065	0.000000e+00	0.000000e+00	0.000000
25%	10596.250000	0.207583	0.000000e+00	0.000000e+00	90.000000
50%	20669.000000	0.383856	0.000000e+00	0.000000e+00	99.000000
75%	75610.000000	0.713817	1.500000e+07	2.400000e+07	111.000000
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000

	vote_count	vote_average	release_year	budget_adj	revenue_adj
count	10866.000000	10866.000000	10866.000000	1.086600e+04	1.086600e+04
mean	217.389748	5.974922	2001.322658	1.755104e+07	5.136436e+07
std	575.619058	0.935142	12.812941	3.430616e+07	1.446325e+08
min	10.000000	1.500000	1960.000000	0.000000e+00	0.000000e+00
25%	17.000000	5.400000	1995.000000	0.000000e+00	0.000000e+00
50%	38.000000	6.000000	2006.000000	0.000000e+00	0.000000e+00
75%	145.750000	6.600000	2011.000000	2.085325e+07	3.369710e+07
max	9767.000000	9.200000	2015.000000	4.250000e+08	2.827124e+09

1.1.4 Zero Values

The 25% and 50% values in "budget_adj" and "revenue_adj" columns are zero. So we treat these zero values as missing values and will drop their rows.

```
In [5]: df.query('revenue_adj == 0').id.count()
```

```
Out [5]: 6016
```

```
In [6]: df.query('budget_adj == 0').id.count()
```

```
Out [6]: 5696
```

We treat these zero values as missing values, and drop their rows.

1.2 Data Cleaning

1.2.1 Dropping Columns

First we drop the columns we discussed earlier.

```
In [7]: df.drop(['imdb_id', 'budget', 'revenue', 'original_title', 'cast', 'homepage', 'direct  
         'keywords', 'overview', 'production_companies', 'release_date'], axis=1, inplace=True)
```

1.2.2 Check for Duplicates

```
In [8]: sum(df.duplicated())
```

```
Out[8]: 1
```

```
In [9]: df.drop_duplicates(inplace=True)
```

```
In [10]: sum(df.duplicated())
```

```
Out[10]: 0
```

1.2.3 Dropping Missing Values

```
In [11]: df.isnull().sum()
```

```
Out[11]: id                0  
         popularity        0  
         runtime           0  
         genres            23  
         vote_count         0  
         vote_average       0  
         release_year       0  
         budget_adj         0  
         revenue_adj        0  
         dtype: int64
```

```
In [12]: df.dropna(how='any', inplace=True)
```

```
In [13]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 10842 entries, 0 to 10865  
Data columns (total 9 columns):  
id                10842 non-null int64  
popularity        10842 non-null float64  
runtime           10842 non-null int64  
genres            10842 non-null object  
vote_count        10842 non-null int64  
vote_average      10842 non-null float64  
release_year      10842 non-null int64  
budget_adj        10842 non-null float64
```

```
revenue_adj      10842 non-null float64
dtypes: float64(4), int64(4), object(1)
memory usage: 847.0+ KB
```

As explained earlier, We drop the rows with zero value for "revenue_adj" or "budget_adj".

```
In [14]: df = df[df.revenue_adj != 0]
         df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 4849 entries, 0 to 10848
Data columns (total 9 columns):
id                4849 non-null int64
popularity        4849 non-null float64
runtime           4849 non-null int64
genres            4849 non-null object
vote_count        4849 non-null int64
vote_average      4849 non-null float64
release_year      4849 non-null int64
budget_adj        4849 non-null float64
revenue_adj       4849 non-null float64
dtypes: float64(4), int64(4), object(1)
memory usage: 378.8+ KB
```

```
In [15]: df = df[df.budget_adj != 0]
         df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 3854 entries, 0 to 10848
Data columns (total 9 columns):
id                3854 non-null int64
popularity        3854 non-null float64
runtime           3854 non-null int64
genres            3854 non-null object
vote_count        3854 non-null int64
vote_average      3854 non-null float64
release_year      3854 non-null int64
budget_adj        3854 non-null float64
revenue_adj       3854 non-null float64
dtypes: float64(4), int64(4), object(1)
memory usage: 301.1+ KB
```

Exploratory Data Analysis

1.2.4 Q1- Which genres were most popular between 1996 and 2015?

Since this question asks about the movies that were released between 1996 and 2015, we filter our data based on "released_year" field.

```
In [16]: df_20years = df[df.release_year > 1995]
```

```
In [17]: df_20years.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2884 entries, 0 to 9153
Data columns (total 9 columns):
id                2884 non-null int64
popularity        2884 non-null float64
runtime           2884 non-null int64
genres            2884 non-null object
vote_count        2884 non-null int64
vote_average      2884 non-null float64
release_year      2884 non-null int64
budget_adj        2884 non-null float64
revenue_adj       2884 non-null float64
dtypes: float64(4), int64(4), object(1)
memory usage: 225.3+ KB
```

In our data set, "genres" column contains multiple values separated by pipe (|) character. In order to find all genres that exist in our data set, we'll run the following code:

```
In [18]: all_genres = list(set(df_20years.genres.str.cat(sep="|").split("|")))
all_genres , len( all_genres)
```

```
Out[18]: (['Music',
           'Comedy',
           'Drama',
           'Fantasy',
           'TV Movie',
           'Horror',
           'War',
           'Mystery',
           'Romance',
           'Animation',
           'Crime',
           'Science Fiction',
           'History',
           'Documentary',
           'Western',
           'Thriller',
           'Foreign',
           'Adventure',
```

```

    'Action',
    'Family'],
    20)

```


There are 20 different genres in our data set.

1.2.5 Number of movies released from 1996 to 2015

```

In [19]: plt.figure(figsize=(10, 5))
         sns.set()
         plt.hist(df_20years['release_year'], bins=20)
         plt.xlabel('year')
         plt.ylabel('Number of Movies')
         plt.title('Number of Movie Released Between 1996 to 2015');

```



As we can see, the number of movies in 2015, is almost twice the number of movies in 1996.

1.2.6 Distribution of popularity between 1996 and 2015:

```

In [20]: df_20years.describe()

```

```


Out[20]:

```

	id	popularity	runtime	vote_count	vote_average	\
count	2884.000000	2884.000000	2884.000000	2884.000000	2884.000000	
mean	49999.492025	1.283314	108.489598	601.418169	6.112067	
std	74838.071719	1.601480	18.801966	945.451907	0.779351	
min	12.000000	0.001117	15.000000	10.000000	2.200000	
25%	8487.750000	0.497205	95.000000	85.750000	5.600000	
50%	13412.000000	0.876550	105.000000	253.000000	6.100000	
75%	58169.250000	1.473959	118.000000	688.250000	6.600000	

max	417859.000000	32.985763	338.000000	9767.000000	8.200000
	release_year	budget_adj	revenue_adj		
count	2884.000000	2.884000e+03	2.884000e+03		
mean	2006.661928	4.761033e+07	1.277516e+08		
std	5.434993	4.827910e+07	2.010761e+08		
min	1996.000000	9.693980e-01	2.370705e+00		
25%	2002.000000	1.339331e+07	1.614076e+07		
50%	2007.000000	3.137941e+07	5.783592e+07		
75%	2011.000000	6.691203e+07	1.528228e+08		
max	2015.000000	4.250000e+08	2.827124e+09		

```
In [21]: plt.hist(df_20years['popularity'])
plt.xlabel('Movie Rating')
plt.ylabel('Number of Movies')
plt.title('Distribution of Movie Ratings Between 1996 to 2015');
```



As we can see the popularity value for most of the movies is less than 3, and the average value for popularity is 1.28.

1.2.7 Number of movies released for differnt genres from 1996 to 2015

We know some movies have multiple genres. In order to do any study per genre, we need to split the rows with multiple genres to several rows; one row per each genre. The result is saved in `df_genres` data frame.

```
In [22]: df_genres = (df_20years.drop('genres', axis=1)
    .join
    (
    df_20years.genres
    .str
    .split('|', expand=True)
    .stack()
    .reset_index(drop=True, level=1)
    .rename('genres')
    ))
df_genres.head()
```

```
Out [22]:
```

	id	popularity	runtime	vote_count	vote_average	release_year	\
0	135397	32.985763	124	5562	6.5	2015	
0	135397	32.985763	124	5562	6.5	2015	
0	135397	32.985763	124	5562	6.5	2015	
0	135397	32.985763	124	5562	6.5	2015	
1	76341	28.419936	120	6185	7.1	2015	

		budget_adj	revenue_adj	genres
0	1.379999e+08	1.392446e+09	Action	
0	1.379999e+08	1.392446e+09	Adventure	
0	1.379999e+08	1.392446e+09	Science Fiction	
0	1.379999e+08	1.392446e+09	Thriller	
1	1.379999e+08	3.481613e+08	Action	

To find the number of movies were produced in each gener, we'll use the following command:

```
In [23]: movies_in_genres = (df_genres.groupby('genres', as_index=False).id.count()
    .sort_values(by='id', ascending=False))
```


movies_in_genres

```
Out [23]:
```

	genres	id
6	Drama	1334
3	Comedy	1038
17	Thriller	900
0	Action	782
1	Adventure	534
14	Romance	505
4	Crime	472
15	Science Fiction	359
7	Family	320
11	Horror	309
8	Fantasy	288
13	Mystery	269
2	Animation	169
12	Music	93
10	History	88

18	War	79
5	Documentary	31
19	Western	27
9	Foreign	9
16	TV Movie	1

```
In [24]: locations = list(range(1,21))
heights = movies_in_genres.iloc[:, 1]
lables = movies_in_genres.iloc[:, 0]
plt.subplots(figsize=(15, 5))
plt.bar(locations, heights, tick_label=lables)
plt.title('Number of Movies Released for Different Genres from 1996 to 2015')
plt.xlabel('Genres')
plt.ylabel('Average rating');
plt.xticks(rotation=90);
```



As we can see some genres like 'Tv Movie' and 'Foreign' include a very small number of movies in our data set. The genres with the highest number of production are 'Drama', 'Comedy', 'Thriller', 'Action' and 'Adventure'.

1.2.8 Popularity for different genres from 1996 to 2015

```
In [25]: genre_popularity= (df_genres.groupby('genres', as_index=False)['popularity'].mean()
.sort_values(by='popularity', ascending=False ))
```

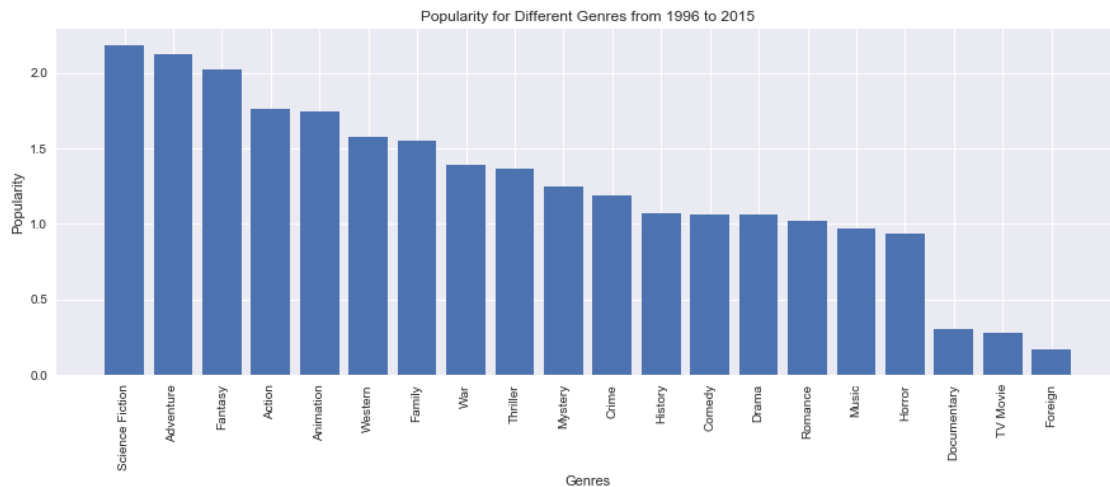
genre_popularity

```
Out [25]:
```

	genres	popularity
15	Science Fiction	2.182806
1	Adventure	2.123756
8	Fantasy	2.025157
0	Action	1.760900

2	Animation	1.744508
19	Western	1.575194
7	Family	1.552059
18	War	1.391474
17	Thriller	1.363990
13	Mystery	1.245874
4	Crime	1.186462
10	History	1.070696
3	Comedy	1.064166
6	Drama	1.063308
14	Romance	1.018031
12	Music	0.970286
11	Horror	0.930450
5	Documentary	0.300557
16	TV Movie	0.273628
9	Foreign	0.164174

```
In [26]: locations = list(range(1,21))
heights = genre_popularity.iloc[:, 1]
lables = genre_popularity.iloc[:, 0]
plt.subplots(figsize=(15, 5))
plt.bar(locations, heights, tick_label=lables)
plt.title('Popularity for Different Genres from 1996 to 2015')
plt.xlabel('Genres')
plt.ylabel('Popularity');
plt.xticks(rotation=90);
```



As we can see 'Science Fiction', 'Adventure', 'Fantasy', 'Action' and 'Animation' had the highest popularity from 1996 to 2015. 'TV Movie' and 'Foreign' genres are the least popular genres and they have the least number of production(we found out in the previous section) in the 20 years of our study.

Now that we've figured out the most popular genres, let's find out the average rating and the trend of popularity for them over these years.

1.2.9 Average rating for the most popular genres

```
In [27]: popular_genres = list(genre_popularity['genres'][:5])
        popular_genres
```

```
Out[27]: ['Science Fiction', 'Adventure', 'Fantasy', 'Action', 'Animation']
```

```
In [28]: df_20years.describe()
```

```
Out[28]:
```

	id	popularity	runtime	vote_count	vote_average \
count	2884.000000	2884.000000	2884.000000	2884.000000	2884.000000
mean	49999.492025	1.283314	108.489598	601.418169	6.112067
std	74838.071719	1.601480	18.801966	945.451907	0.779351
min	12.000000	0.001117	15.000000	10.000000	2.200000
25%	8487.750000	0.497205	95.000000	85.750000	5.600000
50%	13412.000000	0.876550	105.000000	253.000000	6.100000
75%	58169.250000	1.473959	118.000000	688.250000	6.600000
max	417859.000000	32.985763	338.000000	9767.000000	8.200000

	release_year	budget_adj	revenue_adj
count	2884.000000	2.884000e+03	2.884000e+03
mean	2006.661928	4.761033e+07	1.277516e+08
std	5.434993	4.827910e+07	2.010761e+08
min	1996.000000	9.693980e-01	2.370705e+00
25%	2002.000000	1.339331e+07	1.614076e+07
50%	2007.000000	3.137941e+07	5.783592e+07
75%	2011.000000	6.691203e+07	1.528228e+08
max	2015.000000	4.250000e+08	2.827124e+09

```
In [29]: popular_genres_rating = (df_genres.query("genres in @popular_genres")
        .groupby(['genres'], as_index=False)
        .vote_average.mean())

        popular_genres_rating
```

```
Out[29]:
```

	genres	vote_average
0	Action	5.991049
1	Adventure	6.079213
2	Animation	6.335503
3	Fantasy	6.043056
4	Science Fiction	6.002228

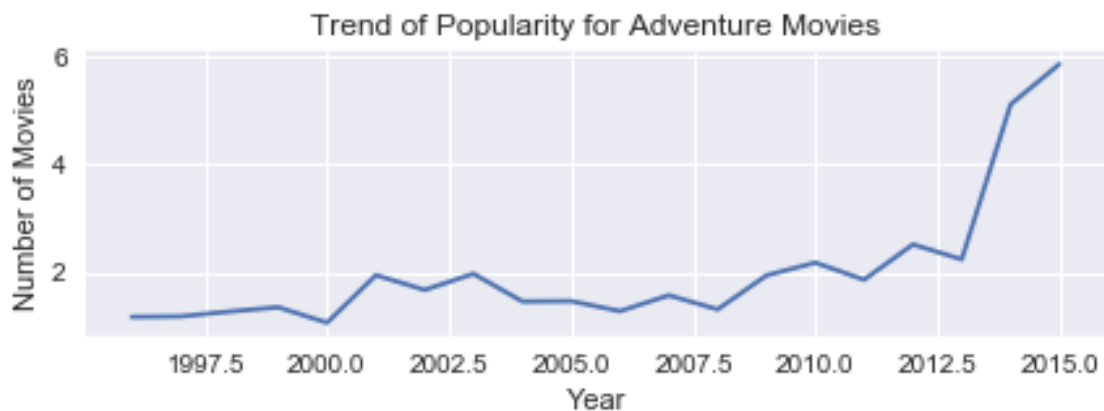

The overall mean value for average rating was '6.11' from 1996 to 2015. Except 'Animation' genre with average rating of '6.33', the other four top popular genres all had a lower value for average voting than the overall mean value.

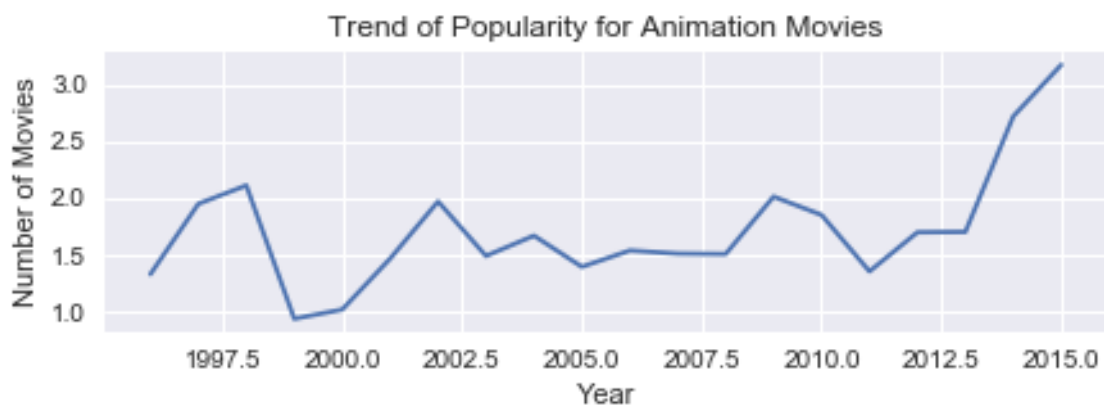


1.2.10 Trend of popularity for most popular genres from 1996 to 2015

```
In [30]: popular_genres_movies = (df_genres.query("genres in @popular_genres")
                                     .groupby(['genres', 'release_year'], as_index=False)
                                     .popularity.mean())
```

Let's see the results in graphs.

```
In [31]: for i, genre in enumerate(popular_genres):
           df1 = popular_genres_movies[popular_genres_movies.genres == genre]
           plt.figure(figsize=(7, 2))
           plt.title('Trend of Popularity for {} Movies'.format(genre))
           plt.xlabel('Year')
           plt.ylabel('Number of Movies')
           plt.plot(df1['release_year'], df1['popularity'])
           plt.show()
```





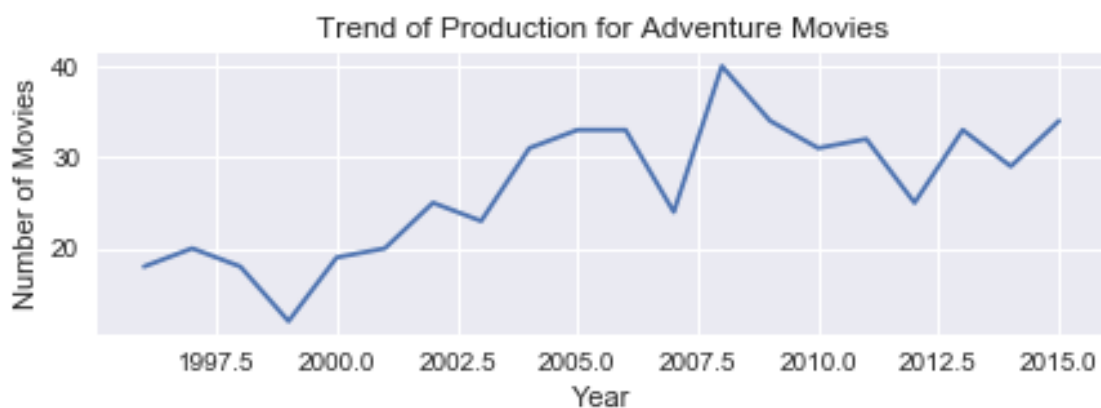

We can see the trend of popularity for the most popular genres has been increased over the 20 years of study. The positive trend is more visible after 2012. The positive trend is more significant for 'Science Fiction', 'Adventure' and 'Action' genres.

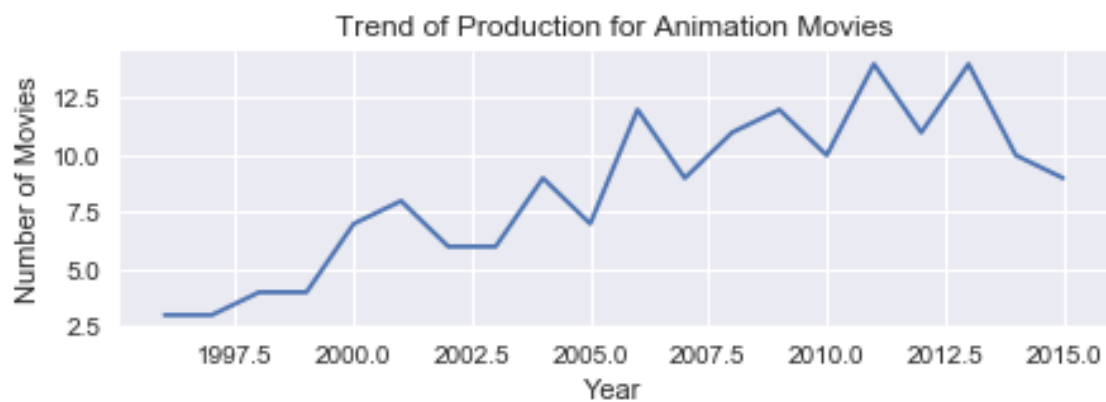
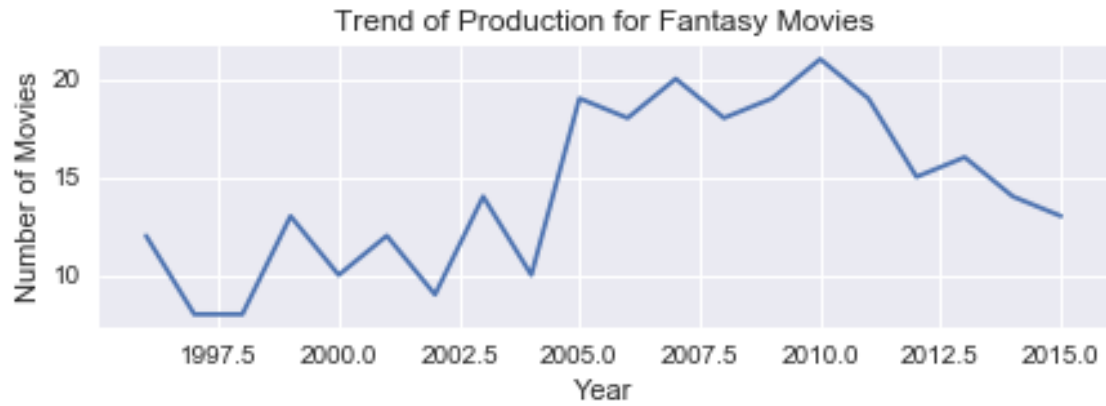
Let's find out how was the trend of movie production for these most popular genres? Did the positive trend in their popularity have any association with their production?

1.2.11 Trend of movie production for most popular genres from 1996 to 2015

```
In [32]: popular_genres_movies_number = (df_genres.query("genres in @popular_genres")
                                             .groupby(['genres', 'release_year'], as_index=False)
                                             .id.count())
```

```
In [33]: for i, genre in enumerate(popular_genres):
           df1 = popular_genres_movies_number[popular_genres_movies_number.genres == genre]
           plt.figure(figsize=(7, 2))
           plt.title('Trend of Production for {} Movies'.format(genre))
           plt.xlabel('Year')
           plt.ylabel('Number of Movies')
           plt.plot(df1['release_year'], df1['id'])
           plt.show()
```






Although we can see that there is an increase in the production of the most popular genres over 20 years, the trend of their production has been negative or none positive after 2012.

1.2.12 Research Q2 : What kind of properties are associated with movies that have high revenues?

To make it possible to use our findings from first question(Research Q1) and make a better conclusion, we use the same time period in this question too. We want to see how many movies have been released during the time period of our study.

```
In [34]: df1= (df_20years.groupby(['release_year'], as_index=False)['id'].count())
```

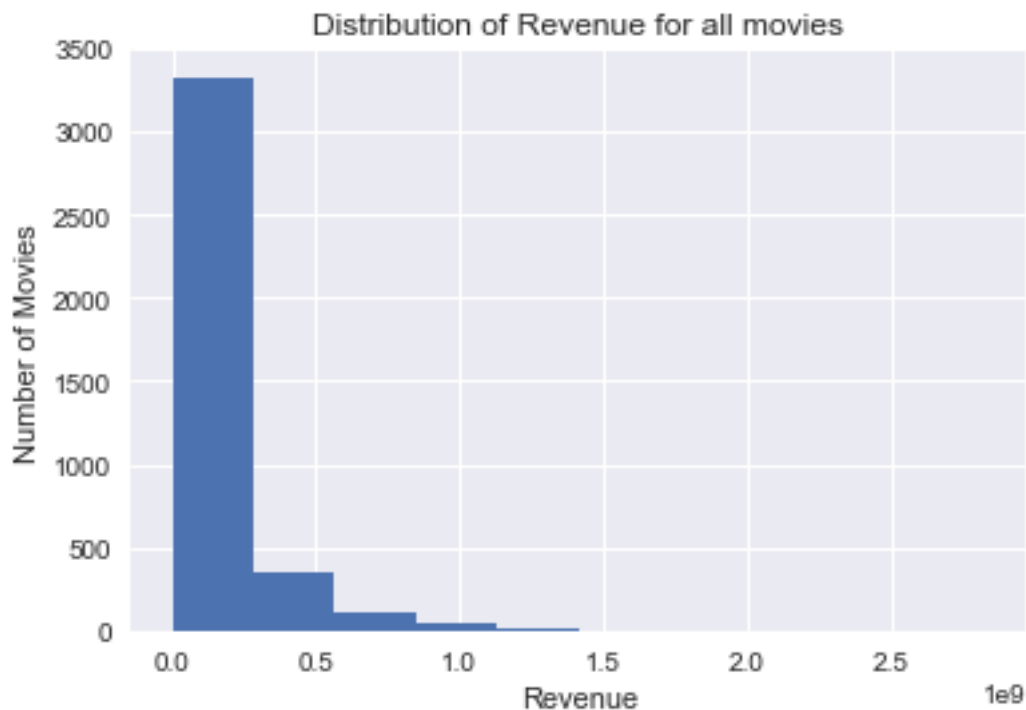
```
sns.set()
locations = list(range(1,21))
heights = df1.iloc[:, 1]
lables = df1.iloc[:, 0]
plt.subplots(figsize=(15, 5))
plt.bar(locations, heights, tick_label=lables)
plt.title('Number of Movies released over years')
plt.xlabel('Years')
plt.ylabel('Number of released movies');
plt.xticks(rotation=90);
```



We can see that number of movies have been increased over the years.

To answer our research question, first we need to find high revenue movies. Let's categorize movies based on some statistics and visualizations.

```
In [35]: sns.set()
plt.hist(df['revenue_adj'])
plt.xlabel('Revenue')
plt.ylabel('Number of Movies')
plt.title('Distribution of Revenue for all movies')
plt.show();
```



```
In [36]: df_20years.describe()
```

```
Out[36]:
```

	id	popularity	runtime	vote_count	vote_average \
count	2884.000000	2884.000000	2884.000000	2884.000000	2884.000000
mean	49999.492025	1.283314	108.489598	601.418169	6.112067
std	74838.071719	1.601480	18.801966	945.451907	0.779351
min	12.000000	0.001117	15.000000	10.000000	2.200000
25%	8487.750000	0.497205	95.000000	85.750000	5.600000
50%	13412.000000	0.876550	105.000000	253.000000	6.100000
75%	58169.250000	1.473959	118.000000	688.250000	6.600000
max	417859.000000	32.985763	338.000000	9767.000000	8.200000

	release_year	budget_adj	revenue_adj
count	2884.000000	2.884000e+03	2.884000e+03
mean	2006.661928	4.761033e+07	1.277516e+08
std	5.434993	4.827910e+07	2.010761e+08
min	1996.000000	9.693980e-01	2.370705e+00
25%	2002.000000	1.339331e+07	1.614076e+07
50%	2007.000000	3.137941e+07	5.783592e+07
75%	2011.000000	6.691203e+07	1.528228e+08
max	2015.000000	4.250000e+08	2.827124e+09

Based on this statistics, we can divide our movies to high_revenue and low_revenue. In order to do that we consider the top 25 percent of revenue as high_revenue movies and the rest as low_revenue.

```
In [37]: df_20years.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2884 entries, 0 to 9153
Data columns (total 9 columns):
id                2884 non-null int64
popularity        2884 non-null float64
runtime           2884 non-null int64
genres            2884 non-null object
vote_count        2884 non-null int64
vote_average      2884 non-null float64
release_year      2884 non-null int64
budget_adj        2884 non-null float64
revenue_adj       2884 non-null float64
dtypes: float64(4), int64(4), object(1)
memory usage: 305.3+ KB
```

```
In [38]: low_revenue = df_20years[df_20years.revenue_adj < 1.528228e+08]
        high_revenue = df_20years[df_20years.revenue_adj > 1.528228e+08]
```

```
In [39]: high_revenue.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 721 entries, 0 to 9043
Data columns (total 9 columns):
id                721 non-null int64
popularity        721 non-null float64
runtime           721 non-null int64
genres            721 non-null object
vote_count        721 non-null int64
vote_average      721 non-null float64
release_year      721 non-null int64
budget_adj        721 non-null float64
revenue_adj       721 non-null float64
dtypes: float64(4), int64(4), object(1)
memory usage: 56.3+ KB
```

```
In [40]: low_revenue.info()
```


```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2163 entries, 12 to 9153
Data columns (total 9 columns):
id                2163 non-null int64
popularity        2163 non-null float64
runtime           2163 non-null int64
genres            2163 non-null object
vote_count        2163 non-null int64
```

```
vote_average    2163 non-null float64
release_year    2163 non-null int64
budget_adj      2163 non-null float64
revenue_adj     2163 non-null float64
dtypes: float64(4), int64(4), object(1)
memory usage: 169.0+ KB
```

```
In [41]: low_revenue_vote = low_revenue['vote_average'].mean()
        high_revenue_vote = high_revenue['vote_average'].mean()
        print(low_revenue_vote, high_revenue_vote)
```

```
6.0293111419324985 6.360332871012483
```

```
In [42]: locations = (1, 2)
        heights = [low_revenue_vote, high_revenue_vote]
        labes = ('Low Revenue', 'High Revenue')
        plt.subplots(figsize=(5, 3))
        plt.bar(locations, heights, tick_label=labes)
        plt.title('Average Rating for High and Low Revenue Movies')
        plt.xlabel('Movies')
        plt.ylabel('Average rating');
```




Average rating for high revenue movies is slightly more than low revenue movies.

```
In [43]: low_revenue_budget = low_revenue['budget_adj'].mean()
        high_revenue_budget = high_revenue['budget_adj'].mean()
        print(low_revenue_budget, high_revenue_budget)
```

30337834.586533476 99427806.69437586

```
In [44]: locations = (1 , 2)
         heights = [low_revenue_budget, high_revenue_budget]
         lables = ('Low Revenue', 'High Revenue')
         plt.subplots(figsize=(5, 3))
         plt.bar(locations, heights, tick_label=lables)
         plt.title('Average Budget for High and Low Revenue Movies')
         plt.xlabel('Movies')
         plt.ylabel('Average Budget');
```



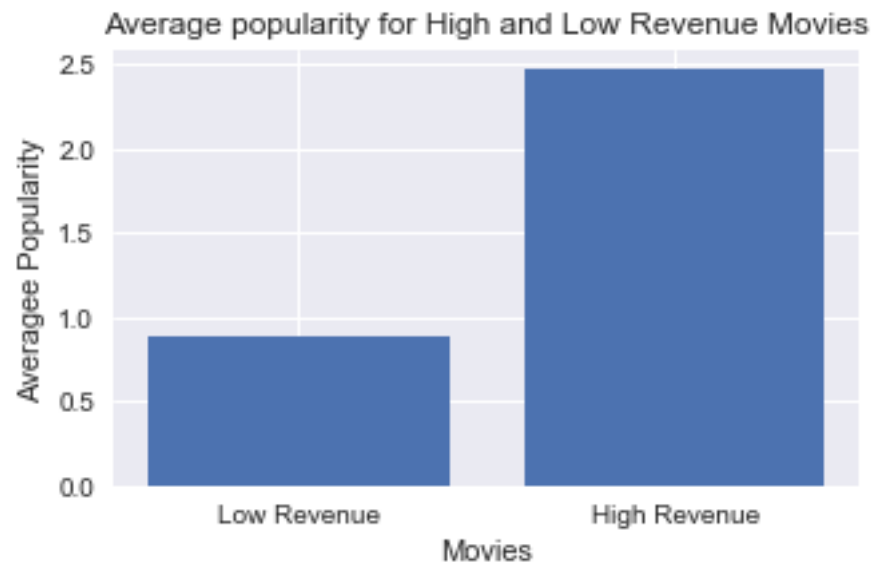
Budget for high revenue movies is 3 times more.

```
In [45]: low_revenue_popularity = low_revenue['popularity'].mean()
         high_revenue_popularity = high_revenue['popularity'].mean()
         print(low_revenue_popularity, high_revenue_popularity)
```

0.8895005903837255 2.4647558141470154

```
In [46]: locations = (1 , 2)
         heights = [low_revenue_popularity, high_revenue_popularity]
         lables = ('Low Revenue', 'High Revenue')
         plt.subplots(figsize=(5, 3))
         plt.bar(locations, heights, tick_label=lables)
         plt.title('Average popularity for High and Low Revenue Movies')
         plt.xlabel('Movies')
         plt.ylabel('Average Popularity')
```

Out[46]: Text(0,0.5,'Averagee Popularity')



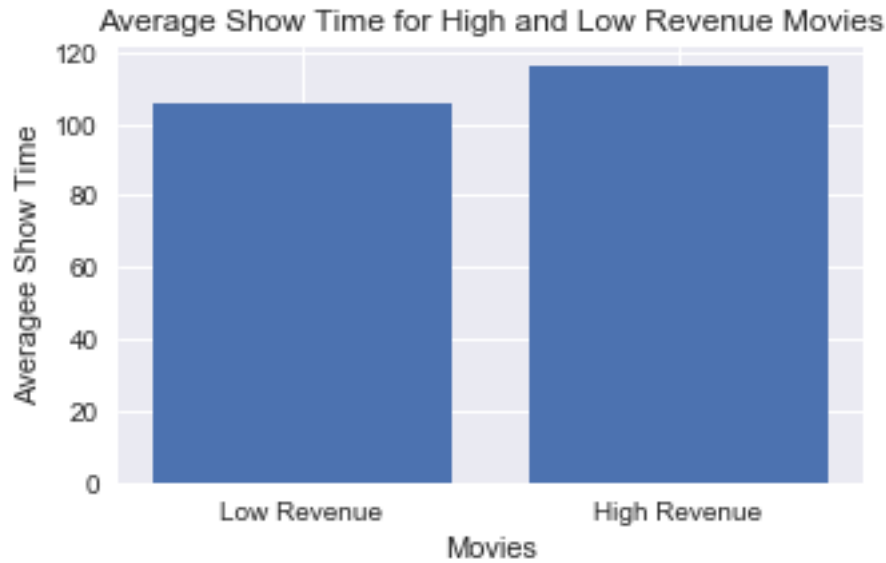
Average popularity for high_revenue movies has been 3 times more than low revenue movies.

```
In [47]: low_revenue_runtime = low_revenue['runtime'].mean()
         high_revenue_runtime = high_revenue['runtime'].mean()
         print(low_revenue_runtime, high_revenue_runtime)
```

106.0610263522885 115.7753120665742

```
In [48]: locations = (1 , 2)
         heights = [low_revenue_runtime, high_revenue_runtime]
         labes = ('Low Revenue', 'High Revenue')
         plt.subplots(figsize=(5, 3))
         plt.bar(locations, heights, tick_label=labes)
         plt.title('Average Show Time for High and Low Revenue Movies')
         plt.xlabel('Movies')
         plt.ylabel('Averagee Show Time')
```

Out[48]: Text(0,0.5,'Averagee Show Time')



The average show time for high revenue movies had been 11 minutes longer than low revenue movies.

```
In [49]: high_revenue_genres = (high_revenue.drop('genres', axis=1)
    .join
    (
    high_revenue.genres
    .str
    .split('|', expand=True)
    .stack()
    .reset_index(drop=True, level=1)
    .rename('genres')
    ))
high_rev_gen_cnt= high_revenue_genres.groupby('genres', as_index = False).id.count()
high_rev_gen_cnt
```

```
Out[49]:
```

	genres	id
0	Action	296
1	Adventure	265
2	Animation	100
3	Comedy	254
4	Crime	99
5	Drama	218
6	Family	152
7	Fantasy	141
8	History	20
9	Horror	41
10	Music	19
11	Mystery	61

12	Romance	107
13	Science Fiction	147
14	Thriller	209
15	War	23
16	Western	7

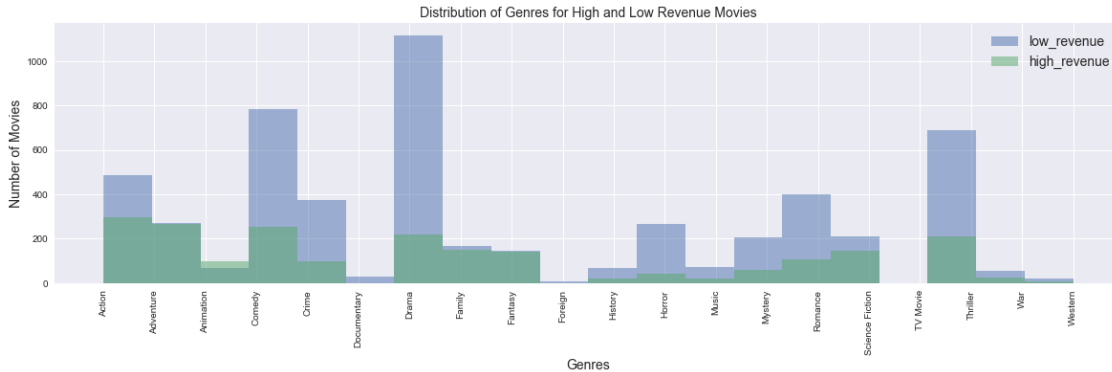
```
In [50]: low_revenue_genres = (low_revenue.drop('genres', axis=1)
    .join
    (
    low_revenue.genres
    .str
    .split('|', expand=True)
    .stack()
    .reset_index(drop=True, level=1)
    .rename('genres')
    ))
low_rev_gen_cnt= low_revenue_genres.groupby('genres', as_index = False).id.count()
low_rev_gen_cnt
```

```
Out [50]:
```

	genres	id
0	Action	486
1	Adventure	269
2	Animation	69
3	Comedy	784
4	Crime	373
5	Documentary	31
6	Drama	1116
7	Family	168
8	Fantasy	147
9	Foreign	9
10	History	68
11	Horror	268
12	Music	74
13	Mystery	208
14	Romance	398
15	Science Fiction	212
16	TV Movie	1
17	Thriller	691
18	War	56
19	Western	20

```
In [51]: sns.set()
plt.figure(figsize=(20, 5))
plt.hist( low_revenue_genres['genres'], label='low_revenue', bins = 20, alpha=0.5)
plt.hist( high_revenue_genres['genres'], label='high_revenue', bins=20, alpha=0.5)
plt.xlabel('Genres', fontsize= 14)
plt.ylabel('Number of Movies', fontsize= 14)
plt.title('Distribution of Genres for High and Low Revenue Movies', fontsize= 14)
```

```
plt.xticks(rotation=90);
plt.legend(fontsize=14);
```



looks like 'Animation' was the most succesful, and 'Documnetory' and 'Foreign' were the lease successful genres in terms of revenue. Most of the movies with 'Animation' genre are categorized as high revenue, and all of the movies in 'Documentary' and 'Foreign' genres are categorized as low revenue. We picked top twenty five percent of revenue as high revenue movies, considering that and the above chart, we can say that 'Adventure', 'Fantasy' and 'Family' were more successful in terms of revenue. 'Action' and 'Science fiction' which were among top popular genres in the first research question(Q1), have a high ranking in revenue too.

Conclusions

In this research, we studied the records of movies from Movie Database(TMDb). Dring the data cleaning, we noticed that more than 60% of our movie records have a zero value in revenue or budget fields. We decided to drop the rows with zero value in budget or revenue fields, so we missed almost 60% of our data.

1.2.13 Findings

- 'Science Fiction', 'Adventure', 'Fantasy', 'Action' and 'Animation' were the most popular genres between 1996 and 2015.
- The trend of popularity has been increasing for the most popular genres over the 20 years of our study. After 2012 the positive trend is more significant.
- The trend of movie production has been positive for the most popular genres between 1996 and 2012, then after 2012 it gets stable or negative.
- There was a positive assosiation between movie revenue and these properties: popularity, average rating, movie showtime and movie budget.
- 'Animation' was the most successful genre in terms of revenue, and the average rating for this genre was higher than the overall mean value.
- Although 'Adventure', 'Fantasy', "Family", 'Action' and 'Sience Fiction' were successful genres in terms of popularity and revenue, their average rating was lower than the overall mean value.

In []: